

## Professional Experience

### Staff Software Engineer

San Francisco, CA

Founding member, Aspire (Python, Typescript, NodeJS, PostgreSQL, Redis)

2018 – 2023

- Joined as the third founding engineer and helped scale the engineering team from 3 to 30 members.
- Architected and implemented the **messaging microservice** facilitating communication for multiple brand users. This service integrated with Gmail, Outlook, and Instagram DM, processing over **millions of messages** daily while maintaining a **99.99% reliability rate**.
- Led the development of a **workflow microservice** which orchestrated and streamlined campaign-critical tasks, helped scale the number of **campaigns 10 fold** and reduced campaign **completion time by 30%**.
- Designed and implemented an **automation-as-a-Service platform**, streamlining custom workflows which reduced **manual interventions by 60%** when collaborating with a creator in a campaign. **Worked closely across teams** to identify and automate key touchpoints.
- Implemented the **shared message inbox** capability in the messaging product which allowed multiple users to access and manage a common email address. This helped the brand delegate and work at scale, resulting in a significant uptick in the volume of communication and ensuring a **50% faster response time**.
- Established a **culture of observability** using Datadog APM throughout the engineering organization. Set and enforced SLOs for every API, ensuring consistent SLA compliance. This led to a reduction in P1 bug resolution time from **10 days to 3 days**.

### Senior Software Engineer

Palo Alto, CA

Adserv, Criteo (Python, DynamoDB, Cassandra, AWS S3)

2016 – 2018

- Achieved annual **cost savings of \$2 million** by meticulously planning and executing a 12-month migration of a recently acquired company's infrastructure serving **100K QPS** from **AWS to proprietary data centers**. Maintaining the **high availability of 99.999%** when migrating live traffic.
- Architected and implemented a **reporting platform** tailored for time-sensitive data, achieving **high-throughput of 40K QPS**. Implemented using worker nodes that execute data rollup tasks and persist results in Cassandra which is later consumed by the reporting service.

### Software Engineer

San Francisco, CA

User mapping system, Brightroll (C, AeroSpike)

2015 – 2016

- Engineered an in-memory, lock-free data structure (Urcu - linux lock-free data structure) for the **user mapping system**—a platform that aligns user identifiers across various systems. This enhancement achieved a **500% reduction in datastore load**.
- Designed and implemented a **cookie-matching server**, adept at synchronizing user identifiers across ad platforms, managing a **60K QPS**. This server streamlines targeted advertising by bridging communication gaps in the ad tech ecosystem.
- Designed a **modular rate-limiting solution** utilizing the token bucket algorithm, enabling granular rate-limiting per partner, per thread. Achieved efficient **rate-limiting performance at 10K QPS**.

### Software Engineer

Bangalore, India

Optimization, Goibibo (Python, Memcached, PostgreSQL)

2010 – 2013

- Engineered an **adaptive expiry algorithm for Memcached**, enhancing cached data management and propelling the **hit rate from 28% to 58%**.
- Developed a **MongoDB-driven monitoring system** for critical logs, designed to issue real-time alerts, resulting in the resolution of **80% of critical logs within the initial hour**.
- Integrated **Pybloomfilter** to optimize API interactions with external vendors, successfully reducing external **API calls by 50%**.

## Programming Languages, Frameworks and Tools

**Programming Languages:** Python, NodeJS, Typescript, React, C

**Databases:** PostgreSQL, Redis, DynamoDB, Google Cloud Task, Google DataStore, Google PubSub

**Cloud platforms:** GCP, AWS

## Education

Masters in Computer Science University Of Southern California, USA	2013 - 2015
B. Tech in Information Technology NIT Allahabad, Allahabad, India	2005 - 2009